

▼ Домашка

tldr:

- Выбрать архитектуру из рассказанных NST, pix2pix, CycleGAN¹
- Подберите к ней задачу, чтобы она вам нравилась
- Подберите еще одну задачу, которая уже решена (если не NST)
- Повторите решение, которое уже есть² (если не NST)
- Решите свою задачу

-
1. Расположены в порядке возрастания сложности и крутизны
 2. Проверьте если вы сделаете этот пункт следующий будет в разы легче

Если вы выбрали Neural Style Transfer

Тут все довольно просто на первый и на второй взгляд. Поэтому недосотаточно просто написать свою функцию потерь и сдать ноутбук. Если вы хотите приличных баллов, то у вас есть две опции:

1. Вы разделяете картинку на две части и переносите на них разные стили.

Нельзя просто взять и два раза применить обычную архитектуру сначала к одной части картинки, а потом к другой.

От вас ожидается, что вы отадите нейросети два картинки стиля и она внутри себя(скорее внутри лосс функции) разделит выходную картинку на две части и к одной части применит один стиль, а к другой - второй.

2. Вы переносите одновременно два стиля на одну картинку контента.

Нельзя просто взять и два раза применить обычную архитектуру сначала с одним стилем, а потом с другим.

От вас ожидается, что вы модифицируете модель(скорее лосс модели) для того, чтобы два стиля учитывались с разными весами.

Если вы выбрали pix2pix

Здесь от вас ожидается, что вы реализуете свою архитектуру для pix2pix модели. Пожалуйста не копируйте код из открытых репозиториев. Этот факт очень легко обнаружить. Перед тем, как приступить проверьте, что обе задачи, которые вы выбрали влезают на вашу видеокарту или на карту Google Colab. Если они не влезают, но вам все равно очень хочется, то вы можете израсходовать все бесплатные триалы облаков(Google, Amazon, .. etc) во вселенной.

Если вы выбрали CycleGAN

Здесь от вас ожидается, что вы реализуете свою архитектуру для CycleGAN модели. Пожалуйста не копируйте код из открытых репозиториев. Этот факт очень легко обнаружить. Перед тем, как приступить проверьте, что обе задачи, которые вы выбрали влезают на вашу видеокарту или на карту Google Colab. CycleGAN в этом смысле хуже, чем pix2pix, он ест больше памяти. Если они влезают, но вам все равно очень хочется, то вы можете израсходовать все бесплатные триалы облаков(Google, Amazon, .. etc) во вселенной.

▼ Remarks:

- Это задание нужно для того, чтобы вы наступили на все грабли, что есть. Узнали об их существовании и научились обходить. Посмотрели на неработающие модели и поняли, что все тлен. Изгуглили весь интернет и в конце заставили это все работать. Проверьте, оно того стиот. Не откладывайте это задание на ночь перед сдачей, так как весь смысл *пуф* улетучится.
- У вас два союзника в этой борьбе:
 1. Оригинальная статья, те психи, что ее писала как то заставили свою модель работать. Их мысли, которыми они спроводили свое детище, позволяют вам написать свой вариант алгоритма.
 2. Гугл, он знает ответы на почти все ваши вопросы, но у него есть две ипостаси одна простая в обещении и вы все ее занаете(русскоязычная), а есть еще одна, которая кусается, но знает больше(англоязычная). Если не знаете языка - уч на ходу :)
- На самом деле у вас есть еще один союзник, это ментор проекта(или лектор или семинарист). Его ресурсом нужно пользоваться в ситуации, в которой вы не можете(занчит попытались и не вышло) найти ответов, используя Гугл и статью.
- Сдавать это все нужно следующим образом. Код вы кидаете на github и отправляете ссылку туда, куда вам сказали(в телеграмм или еще куда-то)

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F

1 from google.colab import drive
2 drive.mount('/content/gdrive/')

↳ Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client\_id=947318989803-6bn6qk8qdgf4n4g3pfj4l4fj3d6c&redirect\_uri=https://colab.research.google.com/notebooks/api
Enter your authorization code:
.....
Mounted at /content/gdrive/

1 !unzip /content/gdrive/My\ Drive/DL\ School/Homeworks/GAN\ and\ Transfer\ Style/maps/train -d train
2 !unzip /content/gdrive/My\ Drive/DL\ School/Homeworks/GAN\ and\ Transfer\ Style/maps/val -d val

1 import pickle
2 import numpy as np
3 from skimage import io
4 from IPython.display import clear_output
5 from tqdm import tqdm, tqdm_notebook
6 from PIL import Image
7 from pathlib import Path
8
9 from torchvision import transforms
10 from multiprocessing.pool import ThreadPool
11 from sklearn.preprocessing import LabelEncoder
12 from torch.utils.data import Dataset, DataLoader
13 import torch.nn as nn
14
15 from matplotlib import colors, pyplot as plt
16 %matplotlib inline
17 import random
18 import cv2
19 # в sklearn не все гладко, чтобы в colab удобно выводить картинки
20 # мы будем игнорировать warnings
21 import warnings
22 warnings.filterwarnings('ignore', category=DeprecationWarning)
```

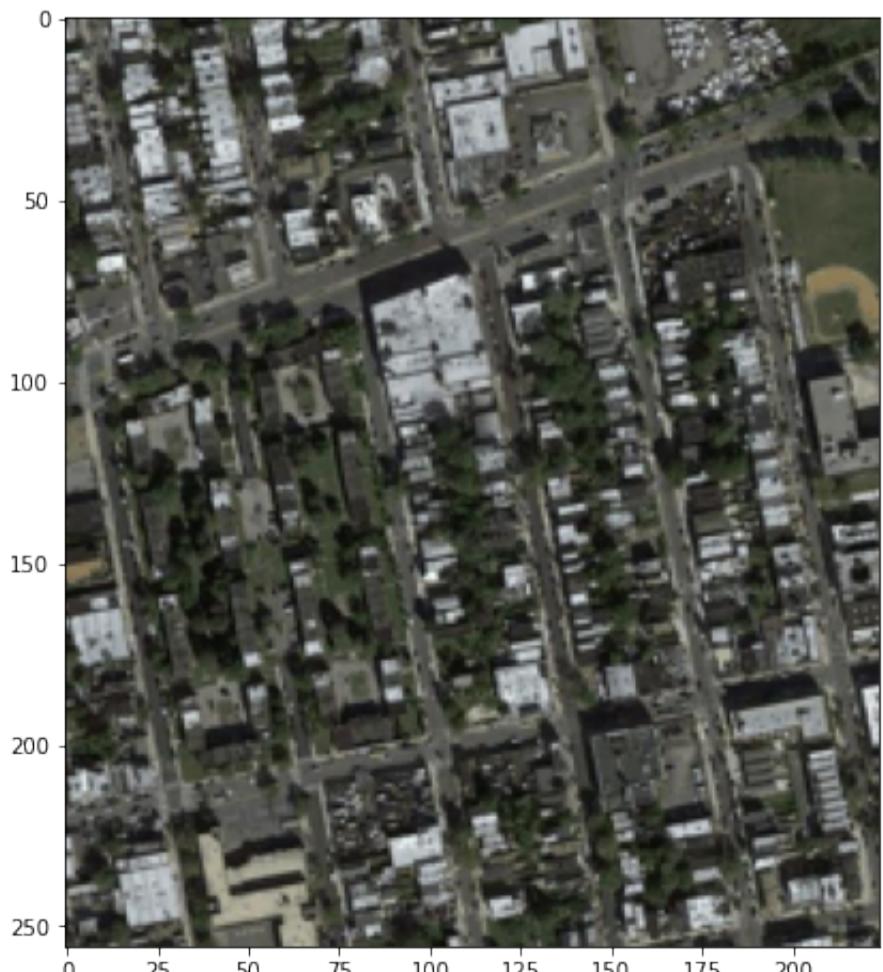
```
1 class ImageDataset(Dataset):
2     def __init__(self, files, mode):
3         super().__init__()
4         # список файлов для загрузки
5         self.files = sorted(files)
6         # режим работы
7         self.mode = mode
8
9         if self.mode not in DATA_MODES:
10             print(f"{self.mode} is not correct; correct modes: {DATA_MODES}")
11             raise NameError
12
13         self.len_ = len(self.files)
14
15     def __len__(self):
16         return self.len_
17
18     def load_sample(self, file):
19         image = Image.open(file)
20         image.load()
21         return image
22
23     def __getitem__(self, index):
24         #add augmentation
25         transform = transforms.Compose([
26             transforms.ToTensor(),
27             transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
28         ])
29         x = self.load_sample(self.files[index])
30         x = self._prepare_sample(x)
31         x = transform(x)
32         return x
33
34     def _prepare_sample(self, image):
35         image = image.resize((RESCALE_SIZE*2, RESCALE_SIZE))
36         return np.array(image)
37
38 # различные режимы датасета
39 DATA_MODES = ['train', 'val']
40 # все изображения будут масштабированы к размеру 224x224 px
41 RESCALE_SIZE = 256
42 # работаем на видеокарте
43 DEVICE = torch.device("cuda")
44
45 def imshow(inp, title=None, plt_ax=plt, default=False):
46     """imshow для тензоров"""
47     inp = inp.numpy().transpose((1, 2, 0))
48     mean = np.array([0.485, 0.456, 0.406])
49     std = np.array([0.229, 0.224, 0.225])
50     inp = std * inp + mean
51     inp = np.clip(inp, 0, 1)
52     plt_ax.imshow(inp)
53     if title is not None:
54         plt_ax.set_title(title)
55     plt_ax.grid(False)
56
57 TRAIN_DIR = Path('train')
58 VAL_DIR = Path('val')
59 train_files = sorted(list(TRAIN_DIR.rglob('*.*jpg')))
60 val_files = sorted(list(VAL_DIR.rglob('*.*jpg')))
61
62 val_dataset = ImageDataset(val_files, mode='val')
63 train_dataset = ImageDataset(train_files, mode='train')
```

```
1 fig, ax = plt.subplots(nrows=3, ncols=3, figsize=(8, 8), \
2                         sharey=True, sharex=True)
3 for fig_x in ax.flatten():
4     random_characters = int(np.random.uniform(0,1000))
5     im_val = val_dataset[random_characters]
6     imshow(im_val[:, :, :224], plt_ax=fig_x)
```



```
1 fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(8, 8), \
2                         sharey=True, sharex=True)
3 random_characters = int(np.random.uniform(0,1000))
4 im_val = val_dataset[0]
5 print(im_val.shape)
6 imshow(im_val[:, :, :224], plt_ax=ax)
7 print(im_val.shape)
```

```
↪ torch.Size([3, 256, 512])
torch.Size([3, 256, 512])
```




```
1 import torchsummary
2 torchsummary.summary(G.cuda(), (3, 256, 256))
```

```
→ -----
      Layer (type)          Output Shape       Param #
=====
      Conv2d-1           [-1, 64, 128, 128]     3,136
      Conv2d-2           [-1, 128, 64, 64]    131,200
      BatchNorm2d-3      [-1, 128, 64, 64]     256
      Conv2d-4           [-1, 256, 32, 32]    524,544
      BatchNorm2d-5      [-1, 256, 32, 32]     512
      Conv2d-6           [-1, 512, 16, 16]   2,097,664
      BatchNorm2d-7      [-1, 512, 16, 16]   1,024
      Conv2d-8           [-1, 512, 8, 8]    4,194,816
      BatchNorm2d-9      [-1, 512, 8, 8]    1,024
      Conv2d-10          [-1, 512, 4, 4]    4,194,816
      BatchNorm2d-11      [-1, 512, 4, 4]    1,024
      Conv2d-12          [-1, 512, 2, 2]    4,194,816
      BatchNorm2d-13      [-1, 512, 2, 2]    1,024
      Conv2d-14          [-1, 512, 1, 1]    4,194,816
      ConvTranspose2d-15  [-1, 512, 2, 2]   4,194,816
      BatchNorm2d-16      [-1, 512, 2, 2]    1,024
      ConvTranspose2d-17  [-1, 512, 4, 4]   8,389,120
      BatchNorm2d-18      [-1, 512, 4, 4]    1,024
      ConvTranspose2d-19  [-1, 512, 8, 8]   8,389,120
      BatchNorm2d-20      [-1, 512, 8, 8]    1,024
      ConvTranspose2d-21  [-1, 512, 16, 16]  8,389,120
      BatchNorm2d-22      [-1, 512, 16, 16]   1,024
      ConvTranspose2d-23  [-1, 256, 32, 32]  4,194,560
      BatchNorm2d-24      [-1, 256, 32, 32]     512
      ConvTranspose2d-25  [-1, 128, 64, 64]  1,048,704
      BatchNorm2d-26      [-1, 128, 64, 64]     256
      ConvTranspose2d-27  [-1, 64, 128, 128]  262,208
      BatchNorm2d-28      [-1, 64, 128, 128]    128
      ConvTranspose2d-29  [-1, 3, 256, 256]   6,147
=====

Total params: 54,419,459
Trainable params: 54,419,459
Non-trainable params: 0

-----
Input size (MB): 0.75
Forward/backward pass size (MB): 54.82
Params size (MB): 207.59
Estimated Total Size (MB): 263.16
```

```
1 torchsummary.summary(D.cuda(), [(3, 256, 256), (3, 256, 256)])
```

```
→ -----
      Layer (type)          Output Shape       Param #
=====
      Conv2d-1           [-1, 64, 128, 128]     6,208
      Conv2d-2           [-1, 128, 64, 64]    131,200
      BatchNorm2d-3      [-1, 128, 64, 64]     256
      Conv2d-4           [-1, 256, 32, 32]    524,544
      BatchNorm2d-5      [-1, 256, 32, 32]     512
      Conv2d-6           [-1, 512, 31, 31]   2,097,664
      BatchNorm2d-7      [-1, 512, 31, 31]   1,024
      Conv2d-8           [-1, 2, 30, 30]    16,386
=====

Total params: 2,777,794
Trainable params: 2,777,794
Non-trainable params: 0

-----
Input size (MB): 147456.00
Forward/backward pass size (MB): 27.52
Params size (MB): 10.60
Estimated Total Size (MB): 147494.12
```

```

1 epochs = 250
2 BCE_loss = nn.BCELoss().to(DEVICE)
3 L1_loss = nn.L1Loss().to(DEVICE)
4 #GAN_loss = nn.DCGANLoss().to(DEVICE)
5 batch_size = 32
6 D_optimizer = torch.optim.Adam(D.parameters(), lr = 2e-4, betas=(0.5, 0.999))
7 G_optimizer = torch.optim.Adam(G.parameters(), lr = 2e-4, betas=(0.5, 0.999))
8 train_hist = {}
9 train_hist['D_losses'] = []
10 train_hist['G_losses'] = []
11 train_hist['per_epoch_ptimes'] = []
12 train_hist['total_ptime'] = []
13 D_scheduler = torch.optim.lr_scheduler.StepLR(D_optimizer, step_size=50, gamma=0.5)
14 G_scheduler = torch.optim.lr_scheduler.StepLR(G_optimizer, step_size=50, gamma=0.5)

1 from torch.autograd import Variable
2 import time
3 G.train()
4 D.train()
5 for i in range(epochs):
6     D_losses = []
7     G_losses = []
8     epoch_start_time = time.time()
9     num_iter = 0
10    train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True, num_workers=4)
11    val_loader = torch.utils.data.DataLoader(val_dataset, batch_size=batch_size, shuffle=True, num_workers=4)
12    for x in train_loader:
13        D.zero_grad()
14        x_ = x[:, :, :, 256:]
15        y_ = x[:, :, :, :256]
16        x_, y_ = Variable(x_.cuda()), Variable(y_.cuda())
17        D_result = D(x_, y_).squeeze()
18        D_real_loss = BCE_loss(D_result, Variable(torch.ones(D_result.size()).cuda()))
19        G_result = G(x_)
20        D_result = D(x_, G_result).squeeze()
21        D_fake_loss = BCE_loss(D_result, Variable(torch.zeros(D_result.size()).cuda()))
22
23        D_train_loss = (D_real_loss + D_fake_loss) * 0.5
24        D_train_loss.backward()
25        D_optimizer.step()
26        train_hist['D_losses'].append(D_train_loss.item())
27        D_losses.append(D_train_loss.item())
28
29        # train generator G
30        G.zero_grad()
31
32        G_result = G(x_)
33        D_result = D(x_, G_result).squeeze()
34        #print(BCE_loss(G_result, y_), BCE_loss(D_result, Variable(torch.ones(D_result.size()).cuda())), L1_loss(G_result, y_))
35        G_train_loss = BCE_loss(D_result, Variable(torch.ones(D_result.size()).cuda())) / 100 + L1_loss(G_result, y_)
36        G_train_loss.backward()
37        G_optimizer.step()
38        train_hist['G_losses'].append(G_train_loss.item())
39
40        G_losses.append(G_train_loss.item())
41        num_iter += 1
42    D_scheduler.step()
43    G_scheduler.step()
44    epoch_end_time = time.time()
45    per_epoch_ptime = epoch_end_time - epoch_start_time
46    print('[%d/%d] - ptime: %.2f, loss_d: %.3f, loss_g: %.3f' % ((i + 1), epochs, per_epoch_ptime, torch.mean(torch.FloatTensor(D_losses)), torch.mean(torch.FloatTensor(G_losses))))

```

↳ [1/250] - ptime: 33.93, loss_d: 0.324, loss_g: 0.636
[2/250] - ptime: 33.81, loss_d: 0.018, loss_g: 0.534
[3/250] - ptime: 33.75, loss_d: 0.007, loss_g: 0.530
[4/250] - ptime: 33.80, loss_d: 0.004, loss_g: 0.531
[5/250] - ptime: 33.61, loss_d: 0.003, loss_g: 0.535
[6/250] - ptime: 33.77, loss_d: 0.002, loss_g: 0.534
[7/250] - ptime: 33.78, loss_d: 0.001, loss_g: 0.537
[8/250] - ptime: 33.68, loss_d: 0.001, loss_g: 0.537

```
[ 9/250] - ptime: 34.09, loss_d: 0.001, loss_g: 0.536
[10/250] - ptime: 33.78, loss_d: 0.001, loss_g: 0.539
[11/250] - ptime: 33.80, loss_d: 0.001, loss_g: 0.538
[12/250] - ptime: 33.77, loss_d: 0.001, loss_g: 0.540
[13/250] - ptime: 33.82, loss_d: 1.101, loss_g: 0.519
[14/250] - ptime: 33.92, loss_d: 0.367, loss_g: 0.478
[15/250] - ptime: 33.56, loss_d: 0.048, loss_g: 0.500
[16/250] - ptime: 33.87, loss_d: 0.083, loss_g: 0.504
[17/250] - ptime: 33.76, loss_d: 0.024, loss_g: 0.512
[18/250] - ptime: 33.75, loss_d: 0.005, loss_g: 0.516
[19/250] - ptime: 33.69, loss_d: 0.005, loss_g: 0.519
[20/250] - ptime: 33.58, loss_d: 0.003, loss_g: 0.521
[21/250] - ptime: 34.06, loss_d: 0.002, loss_g: 0.520
[22/250] - ptime: 33.55, loss_d: 0.002, loss_g: 0.520
[23/250] - ptime: 33.88, loss_d: 0.002, loss_g: 0.519
[24/250] - ptime: 33.89, loss_d: 0.202, loss_g: 0.511
[25/250] - ptime: 33.69, loss_d: 0.227, loss_g: 0.489
[26/250] - ptime: 33.93, loss_d: 0.438, loss_g: 0.496
[27/250] - ptime: 34.02, loss_d: 0.079, loss_g: 0.488
[28/250] - ptime: 33.83, loss_d: 0.013, loss_g: 0.496
[29/250] - ptime: 33.83, loss_d: 0.009, loss_g: 0.500
[30/250] - ptime: 33.69, loss_d: 0.004, loss_g: 0.501
[31/250] - ptime: 33.68, loss_d: 0.003, loss_g: 0.501
[32/250] - ptime: 33.62, loss_d: 0.003, loss_g: 0.503
[33/250] - ptime: 33.68, loss_d: 0.002, loss_g: 0.501
[34/250] - ptime: 33.55, loss_d: 0.002, loss_g: 0.504
[35/250] - ptime: 33.74, loss_d: 0.001, loss_g: 0.500
[36/250] - ptime: 34.39, loss_d: 0.001, loss_g: 0.500
[37/250] - ptime: 34.24, loss_d: 0.001, loss_g: 0.500
[38/250] - ptime: 33.77, loss_d: 0.001, loss_g: 0.500
[39/250] - ptime: 33.90, loss_d: 0.001, loss_g: 0.500
[40/250] - ptime: 33.99, loss_d: 0.001, loss_g: 0.499
[41/250] - ptime: 33.87, loss_d: 0.001, loss_g: 0.500
[42/250] - ptime: 33.96, loss_d: 0.001, loss_g: 0.497
[43/250] - ptime: 34.42, loss_d: 0.001, loss_g: 0.496
[44/250] - ptime: 34.92, loss_d: 0.001, loss_g: 0.495
[45/250] - ptime: 34.71, loss_d: 0.001, loss_g: 0.497
[46/250] - ptime: 34.28, loss_d: 0.001, loss_g: 0.496
[47/250] - ptime: 34.46, loss_d: 0.001, loss_g: 0.494
[48/250] - ptime: 34.27, loss_d: 0.001, loss_g: 0.492
[49/250] - ptime: 34.73, loss_d: 0.001, loss_g: 0.492
[50/250] - ptime: 34.63, loss_d: 0.002, loss_g: 0.492
[51/250] - ptime: 34.74, loss_d: 0.001, loss_g: 0.488
[52/250] - ptime: 35.19, loss_d: 0.000, loss_g: 0.485
[53/250] - ptime: 35.76, loss_d: 0.000, loss_g: 0.484
[54/250] - ptime: 35.68, loss_d: 0.000, loss_g: 0.484
[55/250] - ptime: 35.96, loss_d: 0.000, loss_g: 0.484
[56/250] - ptime: 36.10, loss_d: 0.000, loss_g: 0.484
[57/250] - ptime: 36.32, loss_d: 0.000, loss_g: 0.482
[58/250] - ptime: 35.24, loss_d: 0.000, loss_g: 0.481
[59/250] - ptime: 36.01, loss_d: 0.000, loss_g: 0.481
[60/250] - ptime: 35.82, loss_d: 0.000, loss_g: 0.481
[61/250] - ptime: 35.64, loss_d: 0.000, loss_g: 0.484
[62/250] - ptime: 35.88, loss_d: 0.000, loss_g: 0.481
[63/250] - ptime: 36.08, loss_d: 0.000, loss_g: 0.479
[64/250] - ptime: 35.45, loss_d: 0.000, loss_g: 0.481
[65/250] - ptime: 35.46, loss_d: 0.000, loss_g: 0.478
[66/250] - ptime: 35.22, loss_d: 0.000, loss_g: 0.479
[67/250] - ptime: 35.04, loss_d: 0.000, loss_g: 0.479
[68/250] - ptime: 34.67, loss_d: 0.000, loss_g: 0.478
[69/250] - ptime: 34.90, loss_d: 0.001, loss_g: 0.478
[70/250] - ptime: 34.76, loss_d: 0.001, loss_g: 0.482
[71/250] - ptime: 34.63, loss_d: 0.000, loss_g: 0.479
[72/250] - ptime: 35.06, loss_d: 0.000, loss_g: 0.479
[73/250] - ptime: 34.78, loss_d: 0.000, loss_g: 0.479
[74/250] - ptime: 34.28, loss_d: 0.000, loss_g: 0.479
[75/250] - ptime: 34.29, loss_d: 0.000, loss_g: 0.479
[76/250] - ptime: 34.14, loss_d: 0.000, loss_g: 0.478
[77/250] - ptime: 34.52, loss_d: 0.000, loss_g: 0.478
[78/250] - ptime: 34.26, loss_d: 0.000, loss_g: 0.478
[79/250] - ptime: 34.96, loss_d: 0.000, loss_g: 0.477
[80/250] - ptime: 34.36, loss_d: 0.000, loss_g: 0.477
[81/250] - ptime: 34.43, loss_d: 0.000, loss_g: 0.477
[82/250] - ptime: 34.35, loss_d: 0.000, loss_g: 0.478
[83/250] - ptime: 34.91, loss_d: 0.000, loss_g: 0.477
[84/250] - ptime: 35.47, loss_d: 0.000, loss_g: 0.475
[85/250] - ptime: 34.85, loss_d: 0.000, loss_g: 0.476
[86/250] - ptime: 35.30, loss_d: 0.000, loss_g: 0.475
[87/250] - ptime: 35.49, loss_d: 0.000, loss_g: 0.476
[88/250] - ptime: 35.87, loss_d: 0.000, loss_g: 0.476
```

```
[ 0/250] - ptime: 33.01, loss_d: 0.000, loss_g: 0.470
[ 1/250] - ptime: 36.03, loss_d: 0.000, loss_g: 0.476
[ 2/250] - ptime: 35.45, loss_d: 0.000, loss_g: 0.476
[ 3/250] - ptime: 36.13, loss_d: 0.000, loss_g: 0.476
[ 4/250] - ptime: 35.82, loss_d: 0.000, loss_g: 0.476
[ 5/250] - ptime: 35.75, loss_d: 0.000, loss_g: 0.474
[ 6/250] - ptime: 34.46, loss_d: 0.000, loss_g: 0.474
[ 7/250] - ptime: 34.68, loss_d: 0.000, loss_g: 0.474
[ 8/250] - ptime: 34.06, loss_d: 0.000, loss_g: 0.475
[ 9/250] - ptime: 34.57, loss_d: 0.000, loss_g: 0.473
[10/250] - ptime: 34.34, loss_d: 0.000, loss_g: 0.473
[11/250] - ptime: 34.36, loss_d: 0.000, loss_g: 0.474
[12/250] - ptime: 34.36, loss_d: 0.000, loss_g: 0.473
[13/250] - ptime: 34.23, loss_d: 0.000, loss_g: 0.472
[14/250] - ptime: 34.12, loss_d: 0.000, loss_g: 0.471
[15/250] - ptime: 33.95, loss_d: 0.000, loss_g: 0.471
[16/250] - ptime: 34.41, loss_d: 0.000, loss_g: 0.471
[17/250] - ptime: 34.69, loss_d: 0.000, loss_g: 0.471
[18/250] - ptime: 34.27, loss_d: 0.000, loss_g: 0.470
[19/250] - ptime: 34.32, loss_d: 0.000, loss_g: 0.470
[20/250] - ptime: 34.33, loss_d: 0.000, loss_g: 0.470
[21/250] - ptime: 34.42, loss_d: 0.000, loss_g: 0.471
[22/250] - ptime: 34.31, loss_d: 0.000, loss_g: 0.470
[23/250] - ptime: 34.31, loss_d: 0.000, loss_g: 0.470
[24/250] - ptime: 34.37, loss_d: 0.000, loss_g: 0.470
[25/250] - ptime: 34.56, loss_d: 0.000, loss_g: 0.469
[26/250] - ptime: 34.53, loss_d: 0.000, loss_g: 0.470
[27/250] - ptime: 34.43, loss_d: 0.000, loss_g: 0.469
[28/250] - ptime: 34.38, loss_d: 0.000, loss_g: 0.470
[29/250] - ptime: 34.64, loss_d: 0.000, loss_g: 0.470
[30/250] - ptime: 34.25, loss_d: 0.000, loss_g: 0.469
[31/250] - ptime: 34.51, loss_d: 0.000, loss_g: 0.469
[32/250] - ptime: 33.95, loss_d: 0.000, loss_g: 0.469
[33/250] - ptime: 34.20, loss_d: 0.000, loss_g: 0.468
[34/250] - ptime: 33.93, loss_d: 0.000, loss_g: 0.469
[35/250] - ptime: 33.73, loss_d: 0.000, loss_g: 0.469
[36/250] - ptime: 34.01, loss_d: 0.000, loss_g: 0.469
[37/250] - ptime: 33.96, loss_d: 0.000, loss_g: 0.469
[38/250] - ptime: 34.10, loss_d: 0.000, loss_g: 0.469
[39/250] - ptime: 34.19, loss_d: 0.000, loss_g: 0.468
[40/250] - ptime: 33.85, loss_d: 0.000, loss_g: 0.468
[41/250] - ptime: 33.95, loss_d: 0.000, loss_g: 0.469
[42/250] - ptime: 34.26, loss_d: 0.000, loss_g: 0.468
[43/250] - ptime: 34.06, loss_d: 0.000, loss_g: 0.469
[44/250] - ptime: 34.14, loss_d: 0.000, loss_g: 0.469
[45/250] - ptime: 34.02, loss_d: 0.000, loss_g: 0.469
[46/250] - ptime: 33.99, loss_d: 0.000, loss_g: 0.470
[47/250] - ptime: 33.91, loss_d: 0.000, loss_g: 0.468
[48/250] - ptime: 34.19, loss_d: 0.000, loss_g: 0.469
[49/250] - ptime: 33.53, loss_d: 0.000, loss_g: 0.468
[50/250] - ptime: 33.91, loss_d: 0.000, loss_g: 0.468
[51/250] - ptime: 33.95, loss_d: 0.000, loss_g: 0.469
[52/250] - ptime: 33.65, loss_d: 0.000, loss_g: 0.469
[53/250] - ptime: 33.81, loss_d: 0.000, loss_g: 0.468
[54/250] - ptime: 33.91, loss_d: 0.000, loss_g: 0.468
[55/250] - ptime: 33.96, loss_d: 0.000, loss_g: 0.468
[56/250] - ptime: 33.89, loss_d: 0.000, loss_g: 0.469
[57/250] - ptime: 33.81, loss_d: 0.000, loss_g: 0.468
[58/250] - ptime: 33.96, loss_d: 0.000, loss_g: 0.469
[59/250] - ptime: 33.67, loss_d: 0.000, loss_g: 0.469
[60/250] - ptime: 33.74, loss_d: 0.000, loss_g: 0.469
[61/250] - ptime: 34.05, loss_d: 0.000, loss_g: 0.469
[62/250] - ptime: 33.75, loss_d: 0.000, loss_g: 0.468
[63/250] - ptime: 33.84, loss_d: 0.000, loss_g: 0.467
[64/250] - ptime: 33.78, loss_d: 0.000, loss_g: 0.468
[65/250] - ptime: 33.90, loss_d: 0.000, loss_g: 0.467
[66/250] - ptime: 33.79, loss_d: 0.000, loss_g: 0.467
[67/250] - ptime: 33.86, loss_d: 0.000, loss_g: 0.468
[68/250] - ptime: 33.85, loss_d: 0.000, loss_g: 0.467
[69/250] - ptime: 33.93, loss_d: 0.000, loss_g: 0.467
[70/250] - ptime: 33.81, loss_d: 0.000, loss_g: 0.467
[71/250] - ptime: 33.80, loss_d: 0.000, loss_g: 0.467
[72/250] - ptime: 34.32, loss_d: 0.000, loss_g: 0.467
[73/250] - ptime: 34.43, loss_d: 0.000, loss_g: 0.467
[74/250] - ptime: 34.10, loss_d: 0.000, loss_g: 0.468
[75/250] - ptime: 34.21, loss_d: 0.000, loss_g: 0.467
[76/250] - ptime: 34.07, loss_d: 0.000, loss_g: 0.467
[77/250] - ptime: 34.16, loss_d: 0.000, loss_g: 0.468
[78/250] - ptime: 33.96, loss_d: 0.000, loss_g: 0.468
[79/250] - ptime: 34.13, loss_d: 0.000, loss_g: 0.467
```

```
[168/250] - ptime: 33.95, loss_d: 0.000, loss_g: 0.467
[169/250] - ptime: 33.94, loss_d: 0.000, loss_g: 0.466
[170/250] - ptime: 34.05, loss_d: 0.000, loss_g: 0.466
[171/250] - ptime: 34.12, loss_d: 0.000, loss_g: 0.468
[172/250] - ptime: 33.91, loss_d: 0.000, loss_g: 0.468
[173/250] - ptime: 34.05, loss_d: 0.000, loss_g: 0.468
[174/250] - ptime: 34.00, loss_d: 0.000, loss_g: 0.467
[175/250] - ptime: 34.10, loss_d: 0.000, loss_g: 0.468
[176/250] - ptime: 34.19, loss_d: 0.000, loss_g: 0.468
[177/250] - ptime: 34.22, loss_d: 0.000, loss_g: 0.467
[178/250] - ptime: 34.66, loss_d: 0.000, loss_g: 0.467
[179/250] - ptime: 34.17, loss_d: 0.000, loss_g: 0.467
[180/250] - ptime: 34.06, loss_d: 0.000, loss_g: 0.467
[181/250] - ptime: 34.02, loss_d: 0.000, loss_g: 0.468
[182/250] - ptime: 34.17, loss_d: 0.000, loss_g: 0.467
[183/250] - ptime: 33.86, loss_d: 0.000, loss_g: 0.467
[184/250] - ptime: 34.45, loss_d: 0.000, loss_g: 0.467
[185/250] - ptime: 34.17, loss_d: 0.000, loss_g: 0.467
[186/250] - ptime: 34.04, loss_d: 0.000, loss_g: 0.467
[187/250] - ptime: 33.89, loss_d: 0.000, loss_g: 0.467
[188/250] - ptime: 33.71, loss_d: 0.000, loss_g: 0.467
[189/250] - ptime: 34.02, loss_d: 0.000, loss_g: 0.468
[190/250] - ptime: 34.05, loss_d: 0.000, loss_g: 0.467
[191/250] - ptime: 33.98, loss_d: 0.000, loss_g: 0.468
[192/250] - ptime: 34.19, loss_d: 0.000, loss_g: 0.467
[193/250] - ptime: 33.85, loss_d: 0.000, loss_g: 0.467
[194/250] - ptime: 33.93, loss_d: 0.000, loss_g: 0.468
[195/250] - ptime: 33.79, loss_d: 0.000, loss_g: 0.467
[196/250] - ptime: 33.94, loss_d: 0.000, loss_g: 0.468
[197/250] - ptime: 33.78, loss_d: 0.000, loss_g: 0.468
[198/250] - ptime: 33.86, loss_d: 0.000, loss_g: 0.468
[199/250] - ptime: 34.10, loss_d: 0.000, loss_g: 0.468
[200/250] - ptime: 33.83, loss_d: 0.000, loss_g: 0.468
[201/250] - ptime: 33.98, loss_d: 0.000, loss_g: 0.468
[202/250] - ptime: 33.68, loss_d: 0.000, loss_g: 0.467
[203/250] - ptime: 33.84, loss_d: 0.000, loss_g: 0.468
[204/250] - ptime: 34.05, loss_d: 0.000, loss_g: 0.467
[205/250] - ptime: 33.81, loss_d: 0.000, loss_g: 0.467
[206/250] - ptime: 33.89, loss_d: 0.000, loss_g: 0.467
[207/250] - ptime: 33.94, loss_d: 0.000, loss_g: 0.468
[208/250] - ptime: 33.97, loss_d: 0.000, loss_g: 0.467
[209/250] - ptime: 33.68, loss_d: 0.000, loss_g: 0.467
[210/250] - ptime: 33.89, loss_d: 0.000, loss_g: 0.467
[211/250] - ptime: 33.99, loss_d: 0.000, loss_g: 0.468
[212/250] - ptime: 33.84, loss_d: 0.000, loss_g: 0.468
[213/250] - ptime: 33.56, loss_d: 0.000, loss_g: 0.467
[214/250] - ptime: 33.83, loss_d: 0.000, loss_g: 0.468
[215/250] - ptime: 33.73, loss_d: 0.000, loss_g: 0.468
[216/250] - ptime: 33.75, loss_d: 0.000, loss_g: 0.468
[217/250] - ptime: 33.84, loss_d: 0.000, loss_g: 0.467
[218/250] - ptime: 33.77, loss_d: 0.000, loss_g: 0.468
[219/250] - ptime: 33.63, loss_d: 0.000, loss_g: 0.468
[220/250] - ptime: 33.66, loss_d: 0.000, loss_g: 0.468
[221/250] - ptime: 33.28, loss_d: 0.000, loss_g: 0.468
[222/250] - ptime: 33.60, loss_d: 0.000, loss_g: 0.468
[223/250] - ptime: 33.56, loss_d: 0.000, loss_g: 0.468
[224/250] - ptime: 33.79, loss_d: 0.000, loss_g: 0.468
[225/250] - ptime: 33.88, loss_d: 0.000, loss_g: 0.468
[226/250] - ptime: 33.75, loss_d: 0.000, loss_g: 0.468
[227/250] - ptime: 33.94, loss_d: 0.000, loss_g: 0.469
[228/250] - ptime: 33.45, loss_d: 0.000, loss_g: 0.468
[229/250] - ptime: 33.60, loss_d: 0.000, loss_g: 0.467
[230/250] - ptime: 33.96, loss_d: 0.000, loss_g: 0.468
[231/250] - ptime: 33.14, loss_d: 0.000, loss_g: 0.468
[232/250] - ptime: 34.02, loss_d: 0.000, loss_g: 0.468
[233/250] - ptime: 33.83, loss_d: 0.000, loss_g: 0.468
[234/250] - ptime: 33.91, loss_d: 0.000, loss_g: 0.467
[235/250] - ptime: 33.57, loss_d: 0.000, loss_g: 0.468
[236/250] - ptime: 33.86, loss_d: 0.000, loss_g: 0.468
[237/250] - ptime: 33.86, loss_d: 0.000, loss_g: 0.469
[238/250] - ptime: 33.68, loss_d: 0.000, loss_g: 0.469
[239/250] - ptime: 33.63, loss_d: 0.000, loss_g: 0.468
[240/250] - ptime: 33.85, loss_d: 0.000, loss_g: 0.469
[241/250] - ptime: 33.50, loss_d: 0.000, loss_g: 0.469
[242/250] - ptime: 33.78, loss_d: 0.000, loss_g: 0.469
[243/250] - ptime: 33.58, loss_d: 0.000, loss_g: 0.468
```

```
1 G.eval()
2 D.eval()
```

```
3 val_loader = torch.utils.data.DataLoader(val_dataset, batch_size=batch_size, shuffle=True, num_workers=4)
4 print(len(val_loader))
5 for x in val_loader:
6     x_ = x[:, :, :, :, 256:]
7     y_ = x[:, :, :, :, :256]
8     x_, y_ = Variable(x_.cuda()), Variable(y_.cuda())
9     test_image = G(x_)
10    fig, ax = plt.subplots(1, 3, figsize=(16, 16))
11    imshow(test_image[0].data.cpu(), plt_ax=ax[0])
12    imshow(y_[0].data.cpu(), plt_ax=ax[1])
13    imshow(x_[0].data.cpu(), plt_ax=ax[2])
```

↳ 35

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:10: RuntimeWarning: More than 20 figures have been opened. To prevent this warning, close figures you are not using.

